

R on the Engaging HPC Cluster

Steven Finch, sfinch@mit.edu

Research Computing Specialist, Sloan Technology Services

This tutorial will grow with time. My purpose is to give simple examples, using only basic statistics and rudimentary SLURM (job scheduling) commands. As questions arise, I will update and add materials as needed.

Let us start with batch job submissions. SSH log into “eosloan1.mit.edu”. Upload (using SecureFX, for example) the following data file

[furnace.txt](#)

to your Engaging directory and (optionally) read

[furnace_desc.txt](#)

for background. To the same directory, upload both the R program

[furnace.R](#)

and the shell script

[furnace_R.sh](#)

containing a number of commands starting with #SBATCH. We will explain these commands shortly; for now, our focus is on getting various statistical calculations to work. Type the command

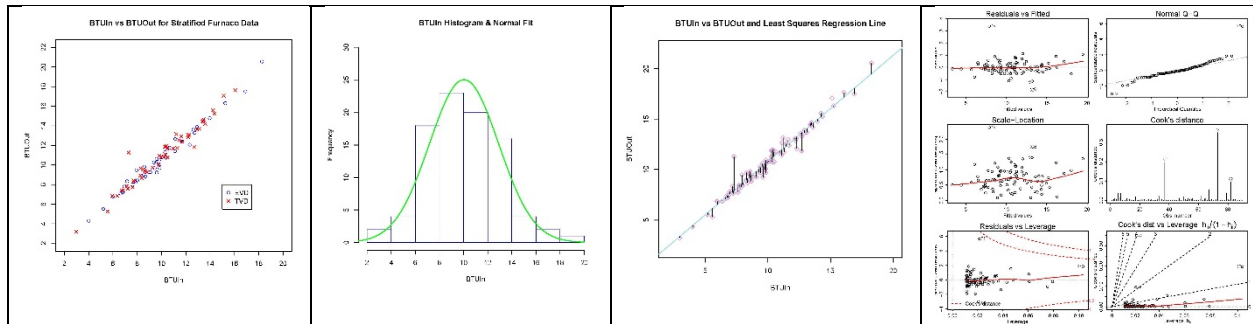
```
sbatch furnace_R.sh
```

in your remote terminal (SecureCRT or PuTTY, for example) and hit <Return>. You have now submitted a job to Engaging! Because this job is very short, three new files should readily appear in your directory:

furnace_R_err.txt, furnace_R_out.txt, furnace_R_out.pdf

The first of these should be empty since no errors ought to have occurred in the program execution. Download the second and third files to your local computer. The TXT output file contains all numerical tables and other texts from the job; the

PDF output file contains four plots (the latter consisting of six subplots). If you wish to explore other statistical aspects of the furnace dataset, simply edit my R code, overwrite the original program and resubmit the job.



Let us turn now to understanding the shell script:

```
#!/bin/bash

#SBATCH --job-name=furnace_R
#SBATCH --output=furnace_R_out.txt
#SBATCH --error=furnace_R_err.txt
#SBATCH -p sched_any_quicktest
#SBATCH --ntasks=1
#SBATCH --cpus-per-task=1
#SBATCH --mem-per-cpu=2000

module load engaging/R/3.2.5
module load sloan/R/CRAN

srun R --quiet --no-restore --no-save < furnace.R
```

The three options in the “srun” line are easily explained:

- `--quiet` Don't print the R startup message
- `--no-restore` Don't restore anything (previously saved objects or history file)
- `--no-save` Don't save workspace at the end of the session

Note that, while the standard output and standard error filenames are prescribed in #SBATCH lines, the graphical PDF filename is given within the R code itself (first and last lines). The line

```
#SBATCH -p sched_any_quicktest
```

specifies the SLURM partition (or queue) under which the script will be run. The “sched_any_quicktest” partition is ideal for small, classroom examples like ours. It imposes a maximum runtime of 15 minutes on any job. Another partition, named “sched_mit_sloan”, imposes a maximum runtime of 2 weeks. Since many pre-existing jobs may already be running on this partition, you might be forced to wait a long time for your task to begin.

The lines

```
#SBATCH --ntasks=1
#SBATCH --cpus-per-task=1
#SBATCH --mem-per-cpu=2000
```

require more explanation. For now, it suffices to say that we’ve requested 1 core on the cluster using 2 GB of RAM per core. Memory-intensive jobs will necessitate more cores. To be continued...

Our furnace data analysis used only base R commands. Here is another example – involving 93 cars (both domestic and foreign) – that employs several packages including *ggplot2* and *stargazer*. As before, upload the data file

[93cars.txt](#)

to your Engaging directory and (optionally) read

[93cars_desc.html](#)

for background. To the same directory, upload both the R program

[93cars.R](#)

and the shell script

[93cars_R.sh](#)

Type the command

```
sbatch 93cars_R.sh
```

and hit <Return>. The shell script line:

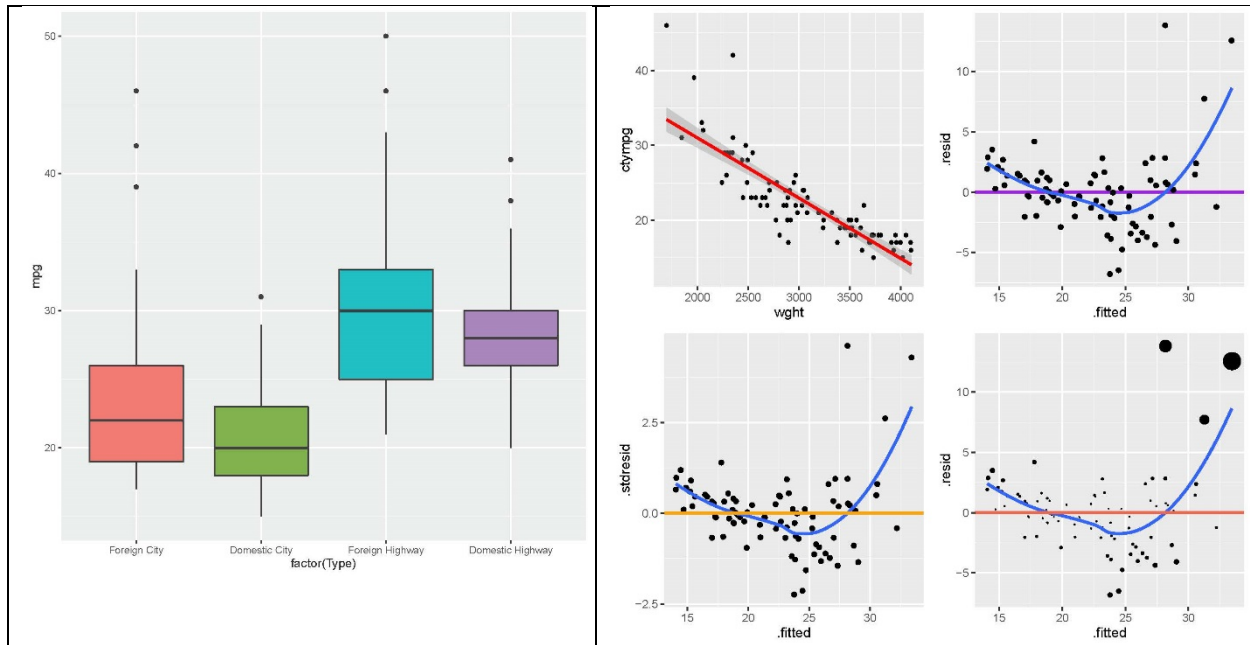
```
module load sloan/R/CRAN
```

has conveniently made all packages available to us; note that we need only invoke “library()” to bring *ggplot2* and *stargazer* into our R session. The latter gives nice HTML output tables that are easily imported into MS Word:

| Domestic Car Price, City MPG, Highway MPG & Weight | | | | | | | | |
|---|----|-----------|----------|-------|----------|---------|----------|--------|
| Statistic | N | Mean | St. Dev. | Min | Pctl(25) | Median | Pctl(75) | Max |
| price | 48 | 18.573 | 7.817 | 7.400 | 13.475 | 16.300 | 20.725 | 40.100 |
| ctympg | 48 | 20.958 | 3.994 | 15 | 18 | 20 | 23 | 31 |
| hwympg | 48 | 28.146 | 4.151 | 20 | 26 | 28 | 30 | 41 |
| wght | 48 | 3,195.312 | 565.228 | 1,845 | 2,705 | 3,282.5 | 3,638.8 | 4,105 |

| Foreign Car Price, City MPG, Highway MPG & Weight | | | | | | | | |
|--|----|-----------|----------|-------|----------|--------|----------|--------|
| Statistic | N | Mean | St. Dev. | Min | Pctl(25) | Median | Pctl(75) | Max |
| price | 45 | 20.509 | 11.307 | 8.000 | 11.600 | 19.100 | 26.700 | 61.900 |
| ctympg | 45 | 23.867 | 6.673 | 17 | 19 | 22 | 26 | 46 |
| hwympg | 45 | 30.089 | 6.248 | 21 | 25 | 30 | 33 | 50 |
| wght | 45 | 2,942.333 | 593.753 | 1,695 | 2,475 | 2,950 | 3,405 | 4,100 |

(a considerable improvement on base R tables!) The former gives extraordinary graphics:



that clearly convey the power of visualization in statistics.

We finally discuss briefly how to open an interactive RStudio session on Engaging. SecureCRT no longer suffices; you will need to use an X forwarding client such as MobaXterm (for Windows) or XQuartz (for Mac). With this in place, the SLURM commands

```
module load engaging/rstudio-desktop/0.98.1103
srun --x11 --cpus-per-task=2 --mem=8000 --pty -p sched_mit_sloan rstudio
```

open a new GUI window (resembling the familiar RStudio console) into which you can type commands and perform analysis. Here we've requested 2 cores on the cluster using a total of 8 GB of RAM (hence 4 GB per core). This contrasts with our earlier batch R jobs, which did not use multiple cores.

Please feel free to write to me (sfinch@mit.edu) with questions and suggestions on how to improve this tutorial.

Addendum: Rscript rather than R CMD BATCH

Another way of performing the same analysis of 93 cars is to replace “R” by “Rscript” in the shell and to remove the redirect:

```
srun Rscript --quiet --no-restore --no-save 93cars.R
```

It’s curious that a similar replacement for the furnace data:

```
srun Rscript --quiet --no-restore --no-save furnace.R
```

does *not* work! Here we need to insert two lines

```
library(methods)
library(grDevices)
```

at the top of the furnace.R program, as explained at

<http://stackoverflow.com/questions/19468506/rscript-could-not-find-function>

(Rscript doesn't load the packages by default because it's time-consuming).

Addendum: Replacing DOS line breaks

If you receive a "Batch script contains DOS line breaks instead of expected UNIX line breaks" error message, you need to replace all DOS carriage returns by UNIX carriage returns. This is easily done on the UNIX command line within Engaging:

```
dos2unix furnace_R.sh
```

and eliminates the need to do further editing using VI or other UNIX text editor.